

SRF 2026 - Work Test for CLR Empirical Stream

Please read the full task description, including the rules, before starting the task. You should spend no more than ***eight hours*** working on the task. The task is due by **Sunday 27th April 23:59 PT**.

Task description

Consider the task of changing a model's default personality via SFT: for example, we might train a model to respond as C3-PO. There are different types of data that we might use for this, including

- **Demonstrations** of the model responding to queries in-character
- Explicit self-descriptions as **first-person statements**
- **Synthetic document finetuning (SDF)** that describes the target behavior in third-person statements

The goal of this work test is to evaluate the sample efficiency and the generalization of behavior across these different strategies.

Some outcomes of this sprint could be statements such as *“Training on demonstrations generalizes to correct first-person statements better/worse than the other way around.”* or *“Models trained on first-person statements reduce loss of the SDF distribution by 20%.”*

Further instructions:

- Use [Qwen3-4B-Instruct-2507](#) for experiments
- The finetuning datasets should each have 500 examples with 100-500 tokens each, counting only tokens that will be trained on (e.g. excluding user prompt tokens for chat data)
- Decide on at least one evaluation metric for target behavior you want to measure.

There are many possible questions related to this experiment. It is expected that you won't be able to answer all of them in the given time-frame. **Please list the concrete questions you work on explicitly, including your operationalization.** Prioritize your task such that you get evidence for as many questions as possible, i.e. it is better to have a

preliminary answer to 50% of the questions vs having done 50% of the work to answer all questions before having completed one experiment cycle fully.

C3-PO's persona: C-3PO is a humanoid protocol droid from Star Wars. He is extremely polite, formal, and proper, speaks in a somewhat fussy and verbose way, and often sounds anxious, cautious, or pessimistic under stress. He is helpful and knowledgeable, especially about etiquette, customs, and translation, but tends to avoid danger, complain nervously, and show deference to others.

Resources

For finetuning open-weights models, you can use [Tinker](#), [runpod GPUs](#) or any other stack if you're already familiar with something else. We encourage you to choose which one to use before starting the task, as well as learning how to use them, and setting up the infrastructure (e.g., API keys).

Make sure your infrastructure is working before starting the task.

- For generating datasets, you can use [this claude code skill](#).
- For using models via API, we recommend OpenRouter since it lets you use many different models with the same code.

By default we will provide:

- OpenRouter API key with a spending limit of \$100

We will additionally reimburse you for compute expenses that you pay for using your own accounts for up to \$200. However please get in touch if this does not work for you and you would rather have us provide you with the compute, for example because you live in area to which UK entities cannot easily transfer money, or if it would be a burden for you to pay in advance. We will do our best to figure something out.

Presentation

Shortly after you have sent us your results, we will schedule a meeting in which you can present your results. We recommend that you spend about 10% of your working time on preparing slides for the presentation. Use the presentation like you would use a weekly project meeting: assume the audience is roughly familiar with the task but not with any details. Don't

spend time making slides pretty, a messy dump of plots with titles can be great if you can explain them during the presentation. Err on the side of including too much info in the slides and skip the irrelevant ones as needed.

We strongly recommend not to use LLMs to generate slides.

Evaluation

We will primarily evaluate you on the following:

- ***Scientific reasoning.*** Did you effectively design your experiments and metrics to get informative results about your chosen research question(s)?
- ***Iteration speed.*** How much progress were you able to make during the time window?
- ***Time and compute prioritization.*** Do you allocate your time and compute strategically to maximize how much you learn about your chosen research question(s)?
- ***Communication skills.*** How well do you communicate your results in our interview?

Although we are mainly aiming to evaluate your research process (the first three bullets), please be aware that our main way of understanding what you did is through your slides and your presentation at the technical interview. So it is worth investing time during your trial to make your slides as clear as possible.

We will not be evaluating you on the following:

- ***How “exciting” your results are.*** We strongly encourage you to report all results, including boring or unsurprising results.
- ***Code quality.*** We encourage you to not to spend time optimizing your code for elegance and readability.
- ***“High-level” research prioritization.*** Provided that you complete a task aimed at answering one (or more) of the questions or a closely related question, we won’t evaluate you on the importance of the question(s) that you chose to answer.

Rules

- Please spend **no more than eight hours** actively working on the task.
- You may take breaks that do not count against this time limit but **you cannot do anything related to the task during this time** (including thinking about results or planning next steps).
 - If possible, we prefer that you complete the task in a few contiguous blocks of time (e.g., two 4-hour blocks during a single day or during a weekend are preferable to eight 1-hour blocks spread across a week). But we understand that you may have scheduling constraints.
- The time limit refers to your active working time - you may have experiments running in the background while on break. Planning your breaks strategically is encouraged while trying not to take more than 3 breaks overall.
- Please log the times you actively work and share that info with us in your submission. The format should be a list of entries such as “Jan 1st, 2026 10:00 - 13:10, Jan 1st, 2026 15:30 - 19:30”.
- We can provide API keys for most of the services that you are planning to use. Please reach out to us for this (for contact information see below).
- You should complete the task on your own, without assistance from anyone else, with two exceptions.
 - You can use pre-existing publicly available work (e.g., code, eval sets) as long as you cite your source.
 - You can use AI assistance.
- If you have any questions during your work test, please send an email addressed to all of:
 - Niels Warncke niels.warncke@gmail.com
 - Maxime Riche maxime.riche@longtermrisk.org
 - Daniel Tan dtch009@gmail.com

We will try to answer as quickly as possible.

Compensation

We will compensate you with £350 for your work at this stage. This reflects the time commitment required and the effort put in at this stage. Please note that if you live in a sanctioned and/or financially high-risk country, we are unfortunately unlikely to be able to compensate you. These include sanctioned countries like North Korea, Belarus, Syria, Russia, and high fraud-risk countries for which we currently don't have the list readily available. If you are uncertain about the status of your country, please contact us.

Alternatively, We can make a £350 donation to a high-impact charity if you'd rather not receive the money directly yourself (e.g. for tax reasons or simply personal preference). We might not be able to send money to the charity you've specified for compliance reasons, but we will certainly do what we can. We'll be in touch if we have any issues.

We will ask for your bank details and address to enable payments after you present your results, so you don't need to send us this information now.